

Reinforcement Learning

Yajur Ahuja
2016121

Q1

classmate
Date _____
Page _____

Assuming finding & not finding have equal prob of $\frac{1}{2}$
(leaching & waiting)

Q1.

Ex 3.4 $p(s', r | s, a) = p(r | s, a, s') \times p(s' | s, a) > 0$.

$s \in \{high, low\}$. $A(s) \in \{h, w\}$. $A(high) \in \{s, w\}$
 $s' \in \{high, low\}$. $A(s') \in \{s, w\}$. $A(low) \in \{s, w, w\}$.

	s	a	s'	r	$p(r s, a, s')$	$p(s' s, a)$	$p(s', r s, a)$
1.	h	s	h	0	$\frac{1}{2}$	$\frac{1}{2}$	$\frac{1}{4}$
2.	h	s	h	1	$\frac{1}{2}$	$\frac{1}{2}$	$\frac{1}{4}$
3.	h	s	l	0	$\frac{1}{2}$	$(1-\frac{1}{2})$	$\frac{1}{2} \cdot (1-\frac{1}{2})$
4.	h	s	l	1	$\frac{1}{2}$	$(1-\frac{1}{2})$	$\frac{1}{2} \cdot (1-\frac{1}{2})$
5.	h	w	h	0	$\frac{1}{2}$	1	$\frac{1}{2}$
6.	h	w	h	1	$\frac{1}{2}$	1	$\frac{1}{2}$
7.	l	s	h	-3	1	$1-\frac{1}{2}$	$1-\frac{1}{2}$
8.	l	s	l	0	$\frac{1}{2}$	$\frac{1}{2}$	$\frac{1}{4}$
9.	l	s	l	1	$\frac{1}{2}$	$\frac{1}{2}$	$\frac{1}{4}$
10.	l	w	l	0	$\frac{1}{2}$	1	$\frac{1}{2}$
11.	l	w	l	1	$\frac{1}{2}$	1	$\frac{1}{2}$
12.	l	r	h	0	1	1	1

but have $p(s', r | s, a) = 0$

Either it is impossible to reach the states from a cert. certain state and action taken i.e. $p(s' | s, a) = 0$
 or if they did receive the state s' , the reward is not what they wanted i.e. $p(r | s, a, s') = 0$.

Q2

Steps:

- 1) Created a probability Matrix : for all $s, s' \rightarrow p_{ss'} = p(s'|s)$ (25 X 25)
- 2) Created an Expected reward Matrix: R (25 X 1)
- 3) γ : discount factor

We know, that the Bellman equation is a set of linear equations

$$V (25 \times 1) = R (25 \times 1) + P (25 \times 25) * V (25 \times 1)$$

$$V = R + \gamma * P * V$$

$$V - \gamma * P * V = R$$

$$(I - \gamma * P) V = R$$

$$V = (I - \gamma * P)^{-1} * R$$

We do the same and obtain the Values in the code

Output Values:

```
[[-16. ,  0.6, -3.3,  5.8, -13.2],  
 [-15.2, -6.9, -5.5, -5.3, -13.7],  
 [-17.1, -10.8, -8.9, -10.2, -16.4],  
 [-20.8, -14.9, -13.1, -14.7, -20.5],  
 [-27.3, -21.6, -19.9, -21.5, -27.1]]
```

Q3

Exercise 15

Q3. Eg (3.8)

Exercise 15 $G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$

Adding C to all Rewards,

$$G_t' = (R_{t+1} + C) + \gamma(R_{t+2} + C) + \dots = \sum_{k=0}^{\infty} \gamma^k (R_{t+k+1} + C)$$

$$G_t' = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1} + \left[\sum_{k=0}^{\infty} \gamma^k \cdot C \right] = V_c$$

$$G_t' = G_t + C \cdot \left(\frac{1}{1-\gamma} \right) \quad \left[G_t' = G_t + V_c \right]$$

$$V_c = \begin{cases} \infty & \gamma = 1 \\ \frac{C}{1-\gamma} & \gamma < 1 \end{cases}$$

$V_{\pi}(s) = \mathbb{E}[G_t | S_t = s]$ (Bellman's equations)

$V_{\pi}'(s) = \mathbb{E}[G_t' | S_t = s] = \mathbb{E}[G_t + V_c | S_t = s]$

$= \mathbb{E}[G_t | S_t = s] + \mathbb{E}[V_c | S_t = s]$

$= \mathbb{E}[G_t | S_t = s] + V_c \text{ (constant)}$

$$V_c = \begin{cases} \infty & \gamma = 1 \\ \frac{C}{1-\gamma} & 0 < \gamma < 1 \\ C & \gamma = 0 \text{ (only for } R_{t+1}) \end{cases}$$

We see that the values for each state are not affected relatively under all policies. Therefore the sign's don't matter in the case where of continuous tasks.

Exercise 16

classmate
 Date _____
 Page _____

Ex 6. $G_t = R_{t+1} + \gamma R_{t+2} + \dots = \sum_{k=t+1}^T \gamma^{k-t-1} R_k$

$G'_t = (R_{t+1} + c) + (R_{t+2} + c) + \dots = \sum_{k=t+1}^T \gamma^{k-t-1} (R_k + c)$

$\Rightarrow G'_t = \sum_{k=t+1}^T \gamma^{k-t-1} R_k + c \sum_{k=t+1}^T \gamma^{k-t-1}$

$= G_t + c [\gamma^0 + \dots + \gamma^{T-t-1}]$

$= G_t + c \left(\frac{1 - \gamma^{T-t}}{1 - \gamma} \right) (= V_c)$

$V'_\pi(s) = \mathbb{E}[G'_t | S_t = s]$

$= \mathbb{E}[G_t + V_c | S_t = s]$

$= \mathbb{E}[G_t | S_t = s] + \mathbb{E}[V_c | S_t = s]$

$= V_\pi(s) + \mathbb{E}\left[\frac{c(1 - \gamma^{T-t})}{1 - \gamma} \middle| S_t = s\right]$

Consider the case of Robot finding the fastest way to reach the terminal point. If the rewards for going ~~longer~~ ^{taking more distance} is ~~to -1~~ ⁻¹, negative, it will for using more distance is -1 for each extra (step/distance) minimizing it will lead to reach the end. But in case of +1 reward, the ~~max~~ minimizing the reward will not lead to the shortest path.

This changes with time step & we see that the Value functions are ~~a~~ relatively ~~diff~~ ^{different}.

Q4 The Bellman's optimality Equation is non-linear. Solved it by using Policy iteration methods given in the code.

The Final Output for both the methods:

Policy Iteration:

Value matrix:

```
[[22.  24.4 22.  19.4 17.5]
 [19.8 22.  19.8 17.8 16. ]
 [17.8 19.8 17.8 16.  14.4]
 [16.  17.8 16.  14.4 13. ]
 [14.4 16.  14.4 13.  11.7]]
```

Policy:

```
['R', 'ULDR', 'L', 'ULDR', 'L', 'UR', 'U', 'UL', 'L', 'L', 'UR', 'U', 'UL', 'UL', 'UL', 'UR', 'U', 'UL', 'UL', 'UL',
 'UR', 'U', 'UL', 'UL', 'UL']
```

Q6 Coded both policy iteration and value iteration and run all to get the output. We get the optimal policy in both the cases.

```
[ ' ', 'L', 'L', 'LD', 'U', 'UL', 'ULDR', 'D', 'U', 'ULDR', 'DR', 'D', 'UR', 'R', 'R', ' ' ]
```

Successive shots in Policy iteration

```
[ [ 0. -13. -19. -21.]
  [-13. -17. -19. -19.]
  [-19. -19. -17. -13.]
  [-21. -19. -13.  0.]]
[ ' ', 'L', 'L', 'L', 'L', 'U', 'UL', 'L', 'D', 'U', 'U', 'D', 'D', 'U', 'R', 'R', ' ' ]
=====
[ [ 0.  0. -1. -2.]
  [ 0. -1. -2. -1.]
  [-1. -2. -1.  0.]
  [-2. -1.  0.  0.]]
[ ' ', 'L', 'L', 'L', 'LD', 'U', 'UL', 'ULDR', 'D', 'U', 'ULDR', 'DR', 'D', 'UR', 'R', 'R', ' ' ]
=====
[ [ 0.  0. -1. -2.]
  [ 0. -1. -2. -1.]
  [-1. -2. -1.  0.]
  [-2. -1.  0.  0.]]
[ ' ', 'L', 'L', 'L', 'LD', 'U', 'UL', 'ULDR', 'D', 'U', 'ULDR', 'DR', 'D', 'UR', 'R', 'R', ' ' ]
=====
```

Successive shots in Value iteration

```
[ [ 0.  0. -1. -1.]
  [ 0. -1. -1. -1.]
  [-1. -1. -1.  0.]
  [-1. -1.  0.  0.]]
[ ' ', 'L', 'L', 'L', 'ULDR', 'U', 'UL', 'ULDR', 'D', 'U', 'ULDR', 'DR', 'D', 'ULDR', 'R', 'R', ' ' ]
=====
[ [ 0.  0. -1. -1.9]
  [ 0. -1. -1.9 -1. ]
  [-1. -1.9 -1.  0. ]
  [-1.9 -1.  0.  0. ]]
[ ' ', 'L', 'L', 'L', 'LD', 'U', 'UL', 'ULDR', 'D', 'U', 'ULDR', 'DR', 'D', 'UR', 'R', 'R', ' ' ]
=====
[ [ 0.  0. -1. -1.9]
  [ 0. -1. -1.9 -1. ]
  [-1. -1.9 -1.  0. ]
  [-1.9 -1.  0.  0. ]]
[ ' ', 'L', 'L', 'L', 'LD', 'U', 'UL', 'ULDR', 'D', 'U', 'ULDR', 'DR', 'D', 'UR', 'R', 'R', ' ' ]
=====
```

We fix the bug by keeping a probability distribution for the actions taken in the policy. As more than one action at a step can have the optimal value, we will divide the probability to each of the actions as (1/ no of actions giving optimal value). Then

we would not have a random selection of action when iterating through the policy iteration algorithm and the algorithm will converge.